


# Quiz 8

---

Create a square polygon data frame with its upper left corner being (180000, 331000), its edge length being 1000, and its attribute being its area. Get the index of the meuse data points that are within the polygon.



Quiz 8 will be due by the end of April 15

# Lecture 20

## Vector/Raster Fusion

---

GEOG 489

SPRING 2020

# Overlay and spatial query

---

**Source layer:** the layer that is used to determine the selection from the target layer based on some spatial relationship.

**Target layer:** the layer to extract data FROM

R function: `over(x, y)`

x: geometry (locations) of the queries

y: layer from which the geometries or attributes are queried

# Overlay and spatial query

over(meuse, srdf)

```
> meuse
```

	coordinates	cadmium	copper	lead	zinc	elev	dist	om	ffreq	soil	lime	landuse	dist.m
1	(181072, 333611)	11.7	85	299	1022	7.909	0.00135803	13.6	1	1	1	Ah	50
2	(181025, 333558)	8.6	81	277	1141	6.983	0.01222430	14.0	1	1	1	Ah	30
3	(181165, 333537)	6.5	68	199	640	7.800	0.10302900	13.0	1	1	1	Ah	150
4	(181298, 333484)	2.6	81	116	257	7.655	0.19009400	8.0	1	2	0	Ga	270
5	(181307, 333330)	2.8	48	117	269	7.480	0.27709000	8.7	1	2	0	Ah	380
6	(181390, 333260)	3.0	61	137	281	7.791	0.36406700	7.8	1	2	0	Ga	470
7	(181165, 333370)	3.2	31	132	346	8.217	0.19009400	9.2	1	2	0	Ah	240
8	(181027, 333363)	2.8	29	150	406	8.490	0.09215160	9.5	1	1	0	Ab	120
9	(181060, 333231)	2.4	37	133	347	8.668	0.18461400	10.6	1	1	0	Ab	240
10	(181232, 333168)	1.6	24	80	183	9.049	0.30970200	6.3	1	2	0	w	420


```
> srdf
```

An object of class "SpatialPolygonsDataFrame"  
Slot "data":

```
      x1 x2
r1    1  5
r2    2  4
r3    3  3
```

```
> over(meuse, srdf)
```

```
      x1 x2
1      1  5
2      1  5
3      1  5
4      1  5
5      1  5
```



The output is, for each point, the data frame of the polygon it intersected.

# Overlay and spatial query

`over(srdf, meuse, fn = mean)`

`> meuse`


	coordinates	cadmium	copper	lead	zinc	elev	dist	om	ffreq	soil	lime	landuse	dist.m
1	(181072, 333611)	11.7	85	299	1022	7.909	0.00135803	13.6	1	1	1	Ah	50
2	(181025, 333558)	8.6	81	277	1141	6.983	0.01222430	14.0	1	1	1	Ah	30
3	(181165, 333537)	6.5	68	199	640	7.800	0.10302900	13.0	1	1	1	Ah	150
4	(181298, 333484)	2.6	81	116	257	7.655	0.19009400	8.0	1	2	0	Ga	270
5	(181307, 333330)	2.8	48	117	269	7.480	0.27709000	8.7	1	2	0	Ah	380
6	(181390, 333260)	3.0	61	137	281	7.791	0.36406700	7.8	1	2	0	Ga	470
7	(181165, 333370)	3.2	31	132	346	8.217	0.19009400	9.2	1	2	0	Ah	240
8	(181027, 333363)	2.8	29	150	406	8.490	0.09215160	9.5	1	1	0	Ab	120
9	(181060, 333231)	2.4	37	133	347	8.668	0.18461400	10.6	1	1	0	Ab	240
10	(181232, 333168)	1.6	24	80	183	9.049	0.30970200	6.3	1	2	0	w	420

`> srdf`

An object of class "SpatialPolygonsDataFrame"

Slot "data":

	x1	x2
r1	1	5
r2	2	4
r3	3	3



`> over(srdf, meuse, fn = mean)`

	cadmium	copper	lead	zinc	elev	dist	om	ffreq	soil	lime	landuse	dist.m
r1	4.036000	44.48000	147.2600	475.8800	8.225760	0.1791637	NA	NA	NA	NA	NA	237.4000
r2	2.910526	40.22807	145.4035	452.0702	8.485649	0.3207399	7.219298	NA	NA	NA	NA	361.2281
r3	2.820833	36.08333	169.1667	484.2500	7.722208	0.2075471	7.350000	NA	NA	NA	NA	261.2500

The output is, for each polygon, the `fn` (mean) value of the data frame of the points falling within it.

# Vector and Raster Fusion


## Extract raster information at vector-defined locations

Raster data: tahoe\_highrez\_brick.tif (3 layers)

```
> tahoe_highrez_brick
class       : RasterBrick
dimensions  : 400, 400, 160000, 3  (nrow, ncol, ncell, nlayers)
resolution  : 5.472863e-06, 5.472863e-06  (x, y)
extent      : -119.9328, -119.9306, 39.28922, 39.29141  (xmin, xmax, ymin, ymax)
coord. ref. : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
data source : Z:\Teaching\GISProgramming\geog489_lecture20_raster_vector_data\taho
names       : tahoe_highrez.1, tahoe_highrez.2, tahoe_highrez.3
min values  :          0,          0,          0
max values  :        255,        255,        255
```

Vector data: tahoe\_highrez\_training\_points.shp

```
> tahoe_highrez_training_points
class       : SpatialPointsDataFrame
features    : 30
extent      : -119.9327, -119.9306, 39.28936, 39.29136  (xmin, xmax, ymin, ymax)
coord. ref. : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
variables   : 2
names       : ID,          SPECIES
min values  : 1, Non-vegetation
max values  : 30,          Tree
```



# Vector and Raster Fusion

## Extract raster information at vector-defined locations

To extract pixel data at specific points:


`extract(x, y):`

x: Raster object

y: Vector object

```
extracted_data <- extract(x=tahoe_highrez_brick,  
  y=tahoe_highrez_training_points)
```

	tahoe_highrez.1	tahoe_highrez.2	tahoe_highrez.3
1	166	47	74
2	165	86	110
3	147	28	41
4	147	55	68
5	119	5	9






# Vector and Raster Fusion

## Extract raster information at vector-defined locations

To link this up with the coordinates of the spatial points

```
extracted_data_w_data <- cbind(extracted_data_w_cellnum,  
tahoe_highrez_training_points@coords))
```

ID	tahoe_highrez.1	tahoe_highrez.2	tahoe_highrez.3	coords.x1	coords.x2
1	166	47	74	-119.9319	39.29078
2	165	86	110	-119.9320	39.29082
3	147	28	41	-119.9320	39.29072
4	147	55	68	-119.9309	39.28994
5	119	5	9	-119.9323	39.28954





# Vector and Raster Fusion

## Extract raster information at vector-defined locations


A point will always return a single value. What about a polygon?

Raster data: tahoe\_highrez\_brick.tif (3 layers)

```
> tahoe_highrez_brick
class      : RasterBrick
dimensions : 400, 400, 160000, 3  (nrow, ncol, ncell, nlayers)
resolution : 5.472863e-06, 5.472863e-06  (x, y)
extent     : -119.9328, -119.9306, 39.28922, 39.29141  (xmin, xmax, ymin, ymax)
coord. ref.: +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
data source: Z:\Teaching\GISProgramming\geog489_lecture20_raster_vector_data\tahoe_
names      : tahoe_highrez.1, tahoe_highrez.2, tahoe_highrez.3
min values  :          0,          0,          0
max values  :         255,         255,         255
```

Vector data: tahoe\_highrez\_training\_polygons.shp

```
> tahoe_highrez_training_polygons
class      : SpatialPolygonsDataFrame
features    : 8
extent     : -119.9326, -119.9311, 39.28956, 39.29135  (xmin, xmax, ymin, ymax)
coord. ref.: +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
variables   : 2
names       : ID,          Class
min values  : 1, Non-vegetated
max values  : 8,          Tree
```



# Vector and Raster Fusion


## Extract raster information at vector-defined locations

A point will always return a single value. What about a polygon?

```
extracted_from_poly <- extract(x=tahoe_highrez_brick,  
                                y=tahoe_highrez_training_polygons)
```

# Every pixel falling within a polygon is extracted  
and stored in a list element.

```
> extracted_from_poly[[1]]  
      tahoe_highrez.1 tahoe_highrez.2 tahoe_highrez.3  
[1,]             31             11             16  
[2,]             65             22             29  
[3,]             65             22             29  
[4,]             78             29             39  
[5,]             86              7             20
```



# Vector and Raster Fusion

## Extract raster information at vector-defined locations


What is the fraction of the pixel covered by a polygon?

```
extracted_from_poly <- extract(x=tahoe_highrez_brick,  
y=tahoe_highrez_training_polygons, weights = TRUE, normalizeWeights = FALSE)
```

# Values of 1.0 mean the entire pixel was covered by polygon. Less than 1.0 mean the pixels were at the edge of the polygon, so they weren't entirely covered.

```
> extracted_from_poly[[1]]
```

	tahoe_highrez.1	tahoe_highrez.2	tahoe_highrez.3	weight
[1,]	16	2	6	0.25
[2,]	31	11	16	0.75
[3,]	75	46	47	0.45
[4,]	41	0	1	0.45
[5,]	65	22	29	1.00



# Vector and Raster Fusion


## Extract raster information at vector-defined locations

A lot of times, we want the mean value per polygon, not the entire polygon.

```
extracted_from_poly <- extract(x=tahoe_highrez_brick,  
                                y=tahoe_highrez_training_polygons,  
                                fun=mean)
```

```
> extracted_from_poly
```

	tahoe_highrez.1	tahoe_highrez.2	tahoe_highrez.3
[1,]	147.26455	55.25397	71.51852
[2,]	139.96503	56.95105	66.35664
[3,]	88.05634	26.50704	33.94366
[4,]	254.99089	89.41686	158.08087
[5,]	212.00000	63.55932	91.06780
[6,]	252.62260	72.94030	111.49254
[7,]	179.44118	237.28922	233.55392
[8,]	194.05856	251.87387	252.10360





# Raster Models

Let's perform a simple linear regression, but using extracted data. We are interested in seeing if Lidar derived vegetation height is linearly related to NDVI:

```
# Calculate Lidar height:
```

```
highest_hit_raster <- raster("tahoe_lidar_highesthit.tif")
```

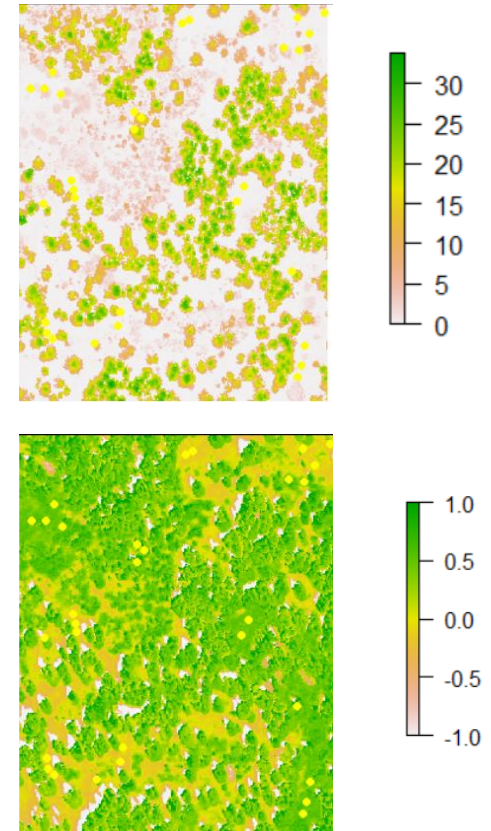
```
bareearth_raster <- raster("tahoe_lidar_bareearth.tif")
```

```
lidar_height <- highest_hit_raster - bareearth_raster
```

```
# Calculate NDVI:
```

```
ndvi <- function(x) { (x[1] - x[2])/(x[1] + x[2]) }
```

```
tahoe_ndvi <- calc(tahoe_highrez_brick, ndvi)
```



# Raster Models

Let's perform a simple linear regression, but using extracted data. We are interested in seeing if Lidar derived vegetation height is linearly related to NDVI:

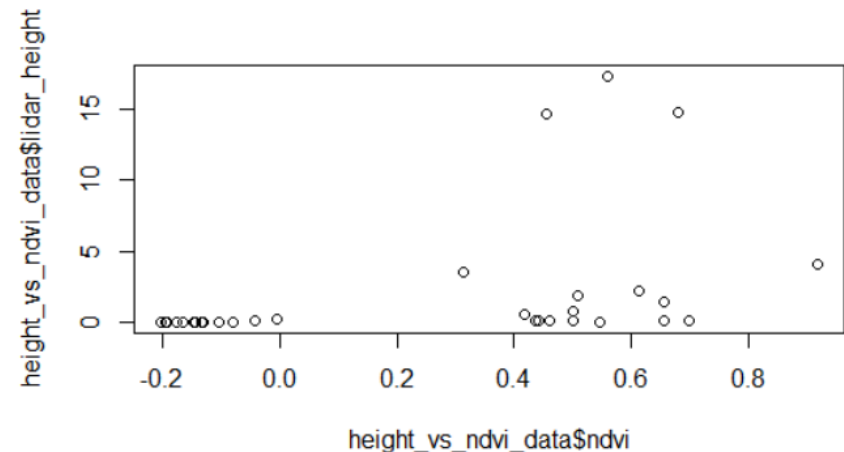
```
# Extract the Lidar height and NDVI for 30 points
```

```
lidar_height_extract <- extract(lidar_height,  
                                tahoe_highrez_training_points,df=TRUE)
```

```
ndvi_extract <- extract(tahoe_ndvi,  
                        tahoe_highrez_training_points,df=TRUE)
```

```
> height_vs_ndvi_data
```

	ID.1	ndvi	ID.2	lidar_height
1	1	0.55868545	1	17.32006836
2	2	0.31474104	2	3.60009766
3	3	0.68000000	3	14.77001953
4	4	0.45544554	4	14.66992188
5	5	0.91935484	5	4.07006836



# Raster Models

---

Let's perform a simple linear regression, but using extracted data. We are interested in seeing if Lidar derived vegetation height is linearly related to NDVI:

```
# Linear regression: response variable ~ explanatory variable (lidar_height ~ ndvi)
```

```
ndvi_height_lm <- lm(height_vs_ndvi_data$lidar_height ~ height_vs_ndvi_data$ndvi)
```

```
Call:
lm(formula = height_vs_ndvi_data$lidar_height ~ height_vs_ndvi_data$ndvi)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-4.2288 -2.6739 -0.3861  0.1059 13.6461
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.8195     0.9867   0.831   0.4132
height_vs_ndvi_data$ndvi  5.1093     2.2535   2.267   0.0313 *
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.395 on 28 degrees of freedom
Multiple R-squared:  0.1551,    Adjusted R-squared:  0.1249
F-statistic:  5.14 on 1 and 28 DF,  p-value: 0.03129
```



# Raster Models

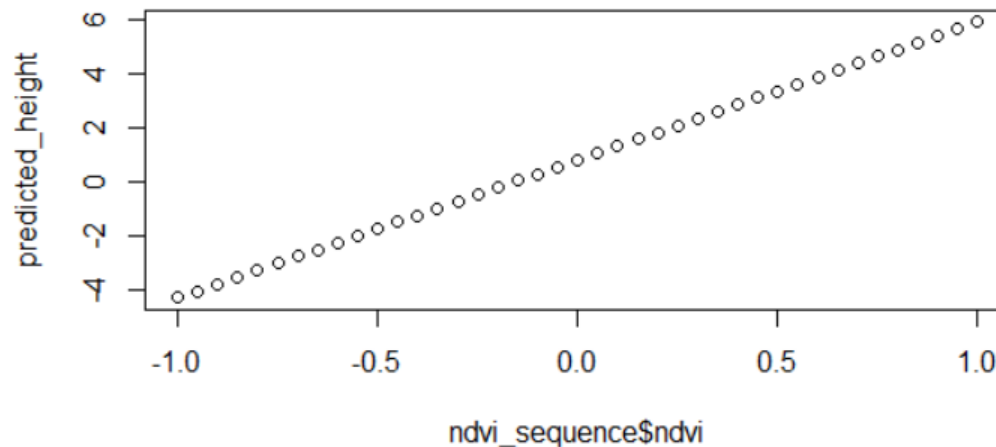
---

Use the linear regression model for PREDICTION

# Let's create a sequence of NDVI values and predict the veg height from the linear regression model:

```
ndvi_sequence = data.frame(ndvi=seq(from=-1, to=1, by=0.05))
```

```
predicted_height <- predict(ndvi_height_lm, ndvi_sequence)
```



# Raster Models

---

Use the linear regression model for PREDICTION


# Let's apply this model to a raster image!

# First off, we need to set the raster layer names to be the same as the variables:

```
names(tahoe_ndvi) <- "ndvi"
```

# Next, we use predict in a similar way:

```
tahoe_height_pred <- predict(tahoe_ndvi, ndvi_height_lm)
```



# Raster Models

---

#What you just learned is INCREDIBLY hard to do in ArcGIS, ENVI, Imagine, etc:

- 1) Extract data from a raster at specific points.
- 2) Create a model from the extracted data.
- 3) Apply the model to a 'test' dataset.
- 4) Apply the model to a raster.

This was TRIVIAL to do in R, and this is one of the chief reasons we use R with GIS!